

Week 2A SQL an...
File

remeber the content

📅 Updated saved memory

已思考 5s >

Got it — I've saved the **Week 2A “SQL and R – Movie Ratings”** assignment requirements (movies/ratings → SQL DB → load into R → handle missing ratings → GitHub submission, no passwords in code, etc.) for future reference.

Week 2A SQL and R Assignment Ov...

Sources



helping me write the code in SQL with the vars in the pdf

已思考 1m 7s >

Here's a clean **starter SQL script** that matches the exact table/column “vars” your PDF suggests:

- users(user_id, name)
- movies(movie_id, title)
- ratings(user_id, movie_id, rating)

Week 2A SQL and R Assignment Ov...

Below are **PostgreSQL (recommended)** and **SQLite** versions, plus a couple **JOIN queries** (also called out in the checklist).

Week 2A SQL and R Assignment Ov...

PostgreSQL version (recommended)

Notes:

- "Uses **PK/FK** and a **junction table** (ratings) like the roadmap describes.
Week 2A SQL and R Assignment Ov... "
- " rating allows decimals (so you can store 3.5). If you want **integers only**, change `NUMERIC(2,1)` → `INTEGER` and the CHECK accordingly."

sql

```
-- =====
-- Week 2A: Movie Ratings DB
-- PostgreSQL
-- =====

DROP TABLE IF EXISTS ratings;
DROP TABLE IF EXISTS movies;
DROP TABLE IF EXISTS users;

-- 1) Dimension tables
CREATE TABLE users (
    user_id INTEGER GENERATED BY DEFAULT AS IDENTITY PRIMARY KEY,
    name    TEXT NOT NULL UNIQUE
);

CREATE TABLE movies (
    movie_id INTEGER GENERATED BY DEFAULT AS IDENTITY PRIMARY KEY,
    title   TEXT NOT NULL UNIQUE
);

-- 2) Junction table (many-to-many)
CREATE TABLE ratings (
    user_id INTEGER NOT NULL REFERENCES users(user_id) ON DELETE CASCADE,
    movie_id INTEGER NOT NULL REFERENCES movies(movie_id) ON DELETE CASCADE,
    rating  NUMERIC(2,1), -- use INTEGER if you want strict 1-5 integers
    PRIMARY KEY (user_id, movie_id),
    CHECK (rating IS NULL OR (rating >= 1 AND rating <= 5))
);
```

```
-- 3) Seed users
INSERT INTO users (name) VALUES
('M'), ('J'), ('W'), ('HP'), ('ZH'), ('CO'), ('ZY');

-- 4) Seed movies/items (edit titles however you want)
INSERT INTO movies (title) VALUES
('Zootopia 2'),
('Stranger Things 5'),
('Captain America 4'),
('Now You See Me 3'),
('Snow White'),
('The SpongeBob Movie: Search for SquarePants');

-- 5) Insert ratings (missing ratings can simply be omitted)
-- This pattern avoids hard-coding ids by joining on name/title.
INSERT INTO ratings (user_id, movie_id, rating)
SELECT u.user_id, m.movie_id, v.rating
FROM (VALUES
('M', 'Zootopia 2', 4),
('M', 'Stranger Things 5', 5),
('M', 'Captain America 4', 3),
('M', 'Now You See Me 3', 4),
('M', 'Snow White', 1),
('M', 'The SpongeBob Movie: Search for SquarePants', 4),

('J', 'Stranger Things 5', 5),
('J', 'The SpongeBob Movie: Search for SquarePants', 4),

('W', 'Zootopia 2', 5),
('W', 'Stranger Things 5', 3),

('HP', 'Captain America 4', 1),

('ZH', 'Now You See Me 3', 4),
('ZH', 'Snow White', 3.5),

('CO', 'Zootopia 2', 4),
('CO', 'The SpongeBob Movie: Search for SquarePants', 3),

('ZY', 'Stranger Things 5', 4),
```

```

('ZY', 'Now You See Me 3', 2)
) AS v(user_name, movie_title, rating)
JOIN users u ON u.name = v.user_name
JOIN movies m ON m.title = v.movie_title;

-- 6) A tidy "analysis-ready" view (this is what you usually pull into R)
SELECT
    u.user_id,
    u.name      AS user_name,
    m.movie_id,
    m.title     AS movie_title,
    r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY u.name, m.title;

```

Optional: make “missing ratings” explicit as NULL

Your checklist mentions verifying missing ratings using `NULL`.

Week 2A SQL and R Assignment Ov...

If you want **every user–movie pair** to exist (rated or not), do this:

sql

```

-- Create a full user x movie grid with NULL ratings (run after seeding users/mov
INSERT INTO ratings (user_id, movie_id, rating)
SELECT u.user_id, m.movie_id, NULL
FROM users u
CROSS JOIN movies m
ON CONFLICT (user_id, movie_id) DO NOTHING;

```

 Copy code

Then you can run:

sql

```

-- Find explicit missing ratings
SELECT u.name, m.title
FROM ratings r
JOIN users u ON u.user_id = r.user_id

```

```
JOIN movies m ON m.movie_id = r.movie_id
WHERE r.rating IS NULL
ORDER BY u.name, m.title;
```

 Copy code

SQLite version (if you're using SQLite)

sql

```
-- =====
-- Week 2A: Movie Ratings DB
-- SQLite
-- =====

PRAGMA foreign_keys = ON;

DROP TABLE IF EXISTS ratings;
DROP TABLE IF EXISTS movies;
DROP TABLE IF EXISTS users;

CREATE TABLE users (
    user_id INTEGER PRIMARY KEY,
    name    TEXT NOT NULL UNIQUE
);

CREATE TABLE movies (
    movie_id INTEGER PRIMARY KEY,
    title   TEXT NOT NULL UNIQUE
);

CREATE TABLE ratings (
    user_id  INTEGER NOT NULL,
    movie_id INTEGER NOT NULL,
    rating   REAL, -- use INTEGER if you want strict integers
    PRIMARY KEY (user_id, movie_id),
    FOREIGN KEY (user_id) REFERENCES users(user_id) ON DELETE CASCADE,
    FOREIGN KEY (movie_id) REFERENCES movies(movie_id) ON DELETE CASCADE,
    CHECK (rating IS NULL OR (rating >= 1 AND rating <= 5))
);
```

```

INSERT INTO users (name) VALUES
('M'), ('J'), ('W'), ('HP'), ('ZH'), ('CO'), ('ZY');

INSERT INTO movies (title) VALUES
('Zootopia 2'),
('Stranger Things 5'),
('Captain America 4'),
('Now You See Me 3'),
('Snow White'),
('The SpongeBob Movie: Search for SquarePants');

INSERT INTO ratings (user_id, movie_id, rating)
SELECT u.user_id, m.movie_id, v.rating
FROM (
VALUES
('M', 'Zootopia 2', 4),
('M', 'Stranger Things 5', 5)
-- ...continue like Postgres version...
) AS v(user_name, movie_title, rating)
JOIN users u ON u.name = v.user_name
JOIN movies m ON m.title = v.movie_title;

```

The 2 queries you'll use most in R

sql

```

-- (1) Tidy joined table (best for importing to R)
SELECT u.name AS user_name, m.title AS movie_title, r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY user_name, movie_title;

-- (2) Per-movie summary (ignores NULL automatically in AVG in Postgres; in SQLIt
SELECT m.title,
       COUNT(r.rating) AS n_ratings,
       AVG(r.rating)   AS avg_rating
FROM movies m
LEFT JOIN ratings r ON r.movie_id = m.movie_id

```

```
GROUP BY m.title
ORDER BY avg_rating DESC;
```

 Copy code

If you tell me **which database you're actually using (Postgres or SQLite)**, I'll tailor it to your exact setup (e.g., identity/serial choice, upsert style, and a "single command" import-friendly `.sql` file layout).

Sources



how to connect to R

To pull your SQL tables into R, the standard approach is **DBI + a database driver**:

- **PostgreSQL** → DBI + RPostgres
- **SQLite** → DBI + RSQlite

Below are copy-paste templates for both 

Option A: PostgreSQL → R (recommended)

1) Install packages

r

 Copy code

```
install.packages(c("DBI", "RPostgres", "dplyr"))
```

2) Connect

r

 Copy code

```
library(DBI)
library(RPostgres)

con <- dbConnect(
  RPostgres::Postgres(),
  host      = "localhost",
  port      = 5432,
  dbname    = "movie_ratings",
  user      = "postgres",
  password = "YOUR_PASSWORD" # better: use env var, see note below
)
```

3) Check tables

r

 Copy code

```
dbListTables(con)
```

4) Pull a joined “analysis-ready” table into R

r

 Copy code

```
sql <- "
SELECT u.name AS user_name, m.title AS movie_title, r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY user_name, movie_title;
"
```

```
ratings_df <- dbGetQuery(con, sql)
ratings_df
```

5) Close connection

r

 Copy code

```
dbDisconnect(con)
```

Credential best practice (no passwords in code):

Use environment variables:

r

 Copy code

```
Sys.getenv("PGPASSWORD")
```

Set once in your `.Renviron` (or Windows env vars), e.g.:

ini

 Copy code

```
PGPASSWORD=your_password_here
```

Then connect without `password=`:

r

 Copy code

```
con <- dbConnect(RPostgres::Postgres(),
                  host="localhost", port=5432,
                  dbname="movie_ratings", user="postgres")
```

Option B: SQLite → R (simplest)

1) Install packages

r

 Copy code

```
install.packages(c("DBI", "RSQLite", "dplyr"))
```

2) Connect to a local `.sqlite` / `.db` file

r

 Copy code

```
library(DBI)
library(RSQLite)
```

```
con <- dbConnect(SQLite(), dbname = "movie_ratings.sqlite")
dbListTables(con)
```

3) Read tables

r

 Copy code

```
users    <- dbReadTable(con, "users")
movies   <- dbReadTable(con, "movies")
ratings <- dbReadTable(con, "ratings")
```

4) Or run a SQL join directly

r

 Copy code

```
ratings_df <- dbGetQuery(con, "
SELECT u.name AS user_name, m.title AS movie_title, r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY user_name, movie_title;
")
```

5) Disconnect

r

 Copy code

```
dbDisconnect(con)
```

Quick R summary (after you've pulled ratings_df)

r

 Copy code

```
library(dplyr)

ratings_df %>%
  group_by(movie_title) %>%
  summarise(
    n_ratings = sum(!is.na(rating)),
    avg_rating = mean(rating, na.rm = TRUE),
    .groups = "drop"
```

```
) %>%  
arrange(desc(avg_rating))
```

If you tell me **which one you're using (Postgres vs SQLite)** and what your **DB name** is (and whether it's local), I'll plug in the exact connection values you should use.



I download the pgadmin4

已思考 10s >

Nice — **pgAdmin 4** is a great start. One important detail:

- **pgAdmin 4 is only a GUI client.**
- To connect from R, you still need a running **PostgreSQL server** (local on your laptop or remote/cloud).

Below is the practical “do this → then this” path.

1) Make sure PostgreSQL server is installed and running

If you installed “PostgreSQL” (not just pgAdmin)

You should have:

- a server running (usually on **localhost**)
- default port **5432**
- a username like **postgres**
- a password you set during installation

If you only installed pgAdmin 4

Then install **PostgreSQL** (the database server).

After that, pgAdmin will connect to the server, and R will connect the same way.

2) In pgAdmin: find the exact connection info R needs

In pgAdmin, right-click your server → **Properties** → **Connection** tab.

Copy these fields:

pgAdmin field	R dbConnect() argument
Host name/address	host
Port	port
Maintenance database (or the DB you created)	dbname
Username	user
Password	password (or environment variable)

Also: create a database like `movie_ratings` in pgAdmin if you haven't:

- Databases → right click → **Create** → **Database...**

3) R: connect to PostgreSQL (recommended way)

Install packages

r

 Copy code

```
install.packages(c("DBI", "RPostgres"))
```

Connect (basic)

r

 Copy code

```
library(DBI)
library(RPostgres)

con <- dbConnect(
```

```
RPostgres::Postgres(),
  host = "localhost",
  port = 5432,
  dbname = "movie_ratings",
  user = "postgres",
  password = "YOUR_PASSWORD"
)

dbGetQuery(con, "SELECT version();") # quick test
dbListTables(con)

dbDisconnect(con)
```

If this works, you're fully connected.

4) Best practice: don't put passwords in your code (for GitHub)

Option A (recommended): put password in .Renvironment

Create/edit the file `~/.Renvironment` and add:

ini

 Copy code

`PGPASSWORD=your_password_here`

ChatGPT 5.2 Thinking

```
library(DBI)
library(RPostgres)

con <- dbConnect(
  RPostgres::Postgres(),
  host = "localhost",
  port = 5432,
  dbname = "movie_ratings",
  user = "postgres"
)
```

Option B (temporary for one session)

r

 Copy code

```
Sys.setenv(PGPASSWORD = "your_password_here")
```

5) Pull your ratings data into R (the join you'll analyze)

r

 Copy code

```
sql <- "
SELECT u.name AS user_name, m.title AS movie_title, r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY user_name, movie_title;
"
```

```
ratings_df <- dbGetQuery(con, sql)
ratings_df
```

If you tell me what you see in pgAdmin:

- Host name/address
- Port
- Username
- Database name you created

...I'll plug them into the exact R `dbConnect()` line for your setup (Windows/local is usually `localhost:5432`).



show me the next steps

Here are the **next steps** from where you are now (pgAdmin 4 installed), in the exact order I'd do them.

1) Confirm PostgreSQL server is actually running

In pgAdmin 4:

1. Open pgAdmin → left panel **Servers**
2. If you see a server (often named **PostgreSQL 16 / PostgreSQL 15**)
 - click it → enter password if asked
3. If you **do NOT** see a server, you likely installed **only pgAdmin** → install **PostgreSQL server** (database engine).

Quick verification inside pgAdmin:

- Open **Query Tool** and run:

sql

```
SELECT version();
```

 Copy code

2) Create your database (one time)

In pgAdmin:

- **Databases** → right click → **Create** → **Database...**
- Name: `movie_ratings` (or any name you want)

3) Run your SQL script to create tables + insert data

In pgAdmin:

1. Click your `movie_ratings` database
2. Open **Tools** → **Query Tool**
3. Paste your SQL (CREATE TABLE + INSERT)
4. Click **Execute**

Confirm tables exist:

- In Query Tool run:

sql

```
SELECT table_name
FROM information_schema.tables
WHERE table_schema='public'
ORDER BY table_name;
```

 Copy code

You should see: movies , ratings , users .

4) In R: install packages

r

```
install.packages(c("DBI", "RPostgres", "dplyr"))
```

 Copy code

5) In R: connect to the same database pgAdmin uses

In pgAdmin, right-click your server → Properties → Connection and copy:

- Host (usually localhost)
- Port (usually 5432)
- Username (often postgres)
- Database name (movie_ratings)

Then in R:

r

```
library(DBI)
library(RPostgres)

con <- dbConnect(
  RPostgres::Postgres(),
  host = "localhost",
  port = 5432,
  dbname = "movie_ratings",
```

```

    user = "postgres",
    password = "YOUR_PASSWORD"
)

dbGetQuery(con, "SELECT version();")
dbListTables(con)

```

 Copy code

6) Pull your “analysis-ready” table into R

r

```

ratings_df <- dbGetQuery(con, "
SELECT u.name AS user_name, m.title AS movie_title, r.rating
FROM ratings r
JOIN users u ON u.user_id = r.user_id
JOIN movies m ON m.movie_id = r.movie_id
ORDER BY user_name, movie_title;
")

```

ratings_df

 Copy code

7) Do your required basic summaries in R

r

```

library(dplyr)

ratings_df %>%
  group_by(movie_title) %>%
  summarise(
    n_ratings = sum(!is.na(rating)),
    avg_rating = mean(rating, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  arrange(desc(avg_rating))

```

 Copy code

8) Handle missing ratings (choose one strategy and show it)

Two common choices:

Strategy A (simple): missing ratings are just absent rows

- In SQL: you **don't insert** a row if someone didn't rate a movie
- In R: you summarize with `na.rm=TRUE` after joins or you keep counts of rated items

Strategy B (explicit NULLs): every user×movie exists, unrated = NULL

- Create a full grid (user × movie) with `NULL` ratings
- Then show `WHERE rating IS NULL` in SQL or `is.na()` in R

(If you want, I'll give you the exact SQL for Strategy B.)

9) Close connection

r

 Copy code

```
dbDisconnect(con)
```

10) GitHub submission checklist

Put in your repo:

- SQL script (create tables + insert)
- R / Quarto file that connects, imports, summarizes, and explains missing strategy
- README explaining how to run it (and **no passwords in code**)

If you want this to be **zero-error**, tell me what pgAdmin shows in **Properties → Connection:**

- Host
- Port
- Username

- Your database name
...and I'll write the exact `dbConnect()` line for your machine.

